

EPITA - Practical Programming



C Programming - 04

Intrusive Structure

Generic Data Structure

Generic in C ?

- using void* equivalence
 - risky (lot of cast)
 - requires separate allocations
- using macro
 - almost unreadable
 - hard to write and maintain

Intrusive Structure ?

Put structures in data:

- One allocation per data
- Data cell can belongs to multiple containers
- Almost transparent implementations

Real life example: list structures in Linux kernel

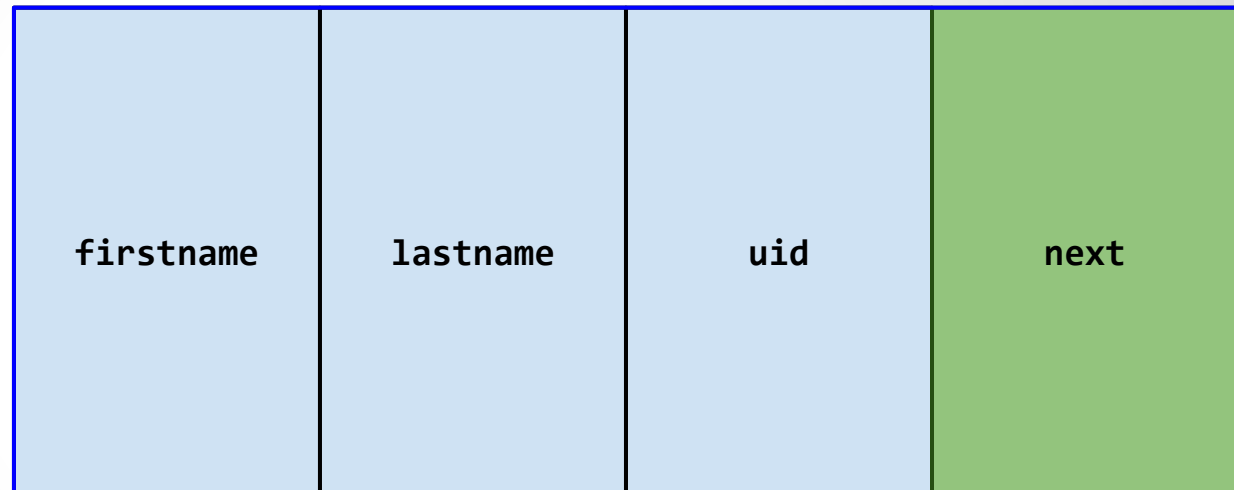
Structures Layout

Structures in Structures

```
struct list {  
    struct list *next;  
};  
  
struct data {  
    char *firstname;  
    char *lastname;  
    unsigned uid;  
    struct list list_;  
};
```

struct data

struct list



list_

Access to parent pointer

- We have a pointer to the `_list` field
- We want the pointer to the outer struct
- We can use `offsetof` from `<stddef.h>`

```
#define CONTAINER_OF(TYPENAME_, FIELDNAME_, PTR_) \
((TYPENAME_*)((char*)PTR_ - offsetof(TYPENAME_, FIELDNAME_)))
```

Intrusive Lists

Intrusive Lists

- Simple list implementation
- External allocation
- Sentinel (not mandatory but useful)
- Few tricks and almost no cast

Intrusive Lists

```
struct list {
    struct list      *next;
};

// Now, compute the parent pointer
# define CONTAINER_OF(TYPENAME_, FIELDNAME_, PTR_) \
    ((TYPENAME_*)((char*)PTR_ - offsetof(TYPENAME_, FIELDNAME_)))

/*
 * list_init(sentinel)
 * Initialize list sentinel
 */
static inline
void list_init(struct list *sentinel) {
    sentinel->next = NULL;
}

/*
 * list_is_empty(l)
 * test for empty list (doesn't work for uninitialized list)
 */
static inline
int list_is_empty(struct list *l) {
    return l->next == NULL;
}
```

Intrusive Lists

```
/*
 * list_push_front(l, cell)
 * add cell in front of the l
 * (keep sentinel unchanged)
 */
static inline
void list_push_front(struct list *l, struct list *cell) {
    cell->next = l->next;
    l->next = cell;
}

/*
 * list_pop_front(l)
 * extract and return the first element of the list
 * returns NULL if the list is empty
 * (keep sentinel unchanged)
 */
static inline
struct list* list_pop_front(struct list *l) {
    struct list *cell = l->next;
    if (cell) {
        l->next = cell->next;
        cell->next = NULL;
    }
    return cell;
}
```

For Each

```
# define foreach(head_, cur) \  
    for (struct list *cur = (head_); cur->next; cur = cur->next)
```

```
size_t len(struct list *head) {  
    size_t    l = 0;  
    foreach(head, it) {  
        l += 1;  
    }  
    return l;  
}
```

For Each

```
struct cell {
    int      value;
    struct list list_;
};

void print_list(struct list *head) {
    foreach(head, it) {
        printf("%d;", CONTAINER_OF(struct cell, list_, it->next)->value);
    }
}
```

Advanced For Each

```
# define foreach(TYPENAME_, FIELD_, head_, elm_) \
    struct list *cur__;\
    for (cur__ = (head_), elm_ = CONTAINER_OF_(TYPENAME_, FIELD_, cur__->next);\
        cur__->next != NULL;\
        cur__ = cur__->next,\
        elm_ = CONTAINER_OF_(TYPENAME_, FIELD_, cur__->next))\

void print_list(struct list *head) {\
    struct cell *elm;\
    foreach(struct cell, list_, head, elm) {\
        printf("%d;", elm->value);\
    }\
}
```

More ?

We Can Do More !

- You can integrate more than one container
- Intrusive approach can be used with any kind of linked containers (like trees)

Other Approach

- Rather than computing offset, we can integrate the parent pointer in the list structure.
- Such implementation can be done using void* based lists.

Even simpler intrusive list

- Another solution is to simply include the list structure as first member
- CONTAINER_OF_ macro is now a simple cast
- With this technique data cell can only appears in one kind of container.

Memory Management

Reference Counting

- Memory management of structured shared by multiple container can become complex
- Life time of memory cell depends on references
- Simple Ref-counting can save a lot of issues.

Reference Counting

```
typedef void (*delete_handle)(void*);

struct refcount {
    size_t count;
    delete_handle delete;
};

static inline
void refcount_init(struct refcount *ref,
delete_handle del) {
    ref->count = 1;
    ref->delete = del;
}
```

```
static inline
void refcount_decr(struct refcount *ref) {
    ref->count -= 1;
    if (ref->count == 0)
        ref->delete(ref);
}

static inline
void refcount_incr(struct refcount *ref) {
    ref->count += 1;
}
```

Usage Example

```
struct data {
    int                value;
    struct list        list_;
    struct refcount    rcount;
};

void delete_data(void *p) {
    free(CONTAINER_OF_(struct data, rcount, p));
}

struct data* create_cell(int x) {
    struct data        *cell;
    cell = malloc(sizeof (struct data));
    cell->value = x;
    refcount_init(&cell->rcount, delete_data);
    return cell;
}
```

Usage Example

```
struct list* build_int_list(int len) {
    struct list      *sentinel;
    sentinel = malloc(sizeof (struct list));
    list_init(sentinel);
    for (int i = 0; i < len; ++i) {
        struct data      *cell;
        cell = create_cell(i);
        // don't need to increment refcount, the only reference is inside the list
        list_push_front(sentinel, &cell->list_);
    }
    return sentinel;
}

void delete_list(struct list *head) {
    while (!list_is_empty(head)) {
        struct data      *cell;
        cell = CONTAINER_OF(struct data, list_, list_pop_front(head));
        // Just decrement refcount
        // If no other ref is available, will free the pointer correctly
        refcount_decr(&cell->rcount);
    }
    free(head); // Sentinel is not ref-counted delete it explicitly
}
```